

Minimal Change to Nominalize the MVP

Ross Tate

Low-Level Systems

int (*)(int)

$[i32] \xrightarrow{A} [i32]$

input/output in
register r5


$[i32] \xrightarrow{B} [i32]$

input/output in
register r3

$[i32] \xrightarrow{C} [i32]$

input/output on
stack

Low-Level Conflicts



Centralizing Coordination

- WebAssembly already requires module centralization within an application
 - All modules needs within an application need to use/import the same exception event
- Exception events are nominal for composability
 - Do not want exception events of different systems to interfere
- Nominal heap types provide composability
 - Representation/casting constraints of different systems do not interfere
- Since application-central module instance is already necessary
 - Use it to coordinate types/RTTs for primitive reference types
- High-level modules are coordinated by the language runtime
 - In WebAssembly, that language runtime is now its own module, the application-central one

Understanding rtt.canon (and rtt.sub)

- `rtt.canon $type : [] -> [(rtt $type)]`
 - Represent `$type` as an automaton
 - Minimize that automaton
 - Canonicalize that minimization
 - Hash that canonicalization
 - Look for preexisting rtt for that type in process-global hashtable
 - If `$type` uses an imported type, do this per instantiation
 - Post-MVP extends with parametric polymorphism
 - Rtt's for type arguments are supplied at run time
 - Above process is done at run time (e.g. per call of polymorphic function constructing pairs in Post-MVP)

Concerns with Structural Typing

- Unknown how to scale to major languages (#116)
 - Whether Post-MVP has practical decidable restriction is longstanding open research problem
 - Post-MVP cast mechanisms cannot express, for example, Java/C# arrays
- Concerns about compile-time performance (#117)
 - Low-level structural types are large, and structural subtyping is quadratic
 - Parametric polymorphism would reduce applicability of memoization
- Issues with composing implementations
 - Unclear if it can be extended with pointer-tagging or unboxed nullary variants (#118)
 - i31ref conflicts with Smis in externref if both are subtypes of anyref (#76/#130)

Observations

- Bottlenecked by dynamic type system
 - Every “this” pointer must be cast
 - Elements of arrays must be cast
- Dynamic type system is nominal
 - RTTs have distinct identities and explicit hierarchy
- Low-level nominal types...
 - have been demonstrated to scale to major languages and existing compilers
 - benefit from fast compile-time and run-time subtyping
 - avoid clashing distinct low-level systems

Suggestion

- Specify hierarchy in types section
 - make heap types generative rather than applicative (i.e. two defs always makes two types)
 - make “sub” relation explicit in type declarations rather than deduced through equi-recursion
 - remove anyref
- Possibly
 - Remove rtt type
 - Separate rtt into GC descriptors (i.e. object headers) and type identifiers (i.e. magic numbers)
 - Enrich GC descriptors to store other meta-data, e.g. v-table